

Universität des Saarlandes
MI Fakultät für Mathematik und Informatik
Department of Computer Science

Bachelorthesis

Capabilities as a Solution against Tracking

submitted by

Tim Christmann
on June 21, 2025

Reviewers

Dr. Sven Bugiel
Noah Mauthe

Eidesstattliche Erklärung

Ich erkläre hiermit an Eides statt, dass ich die vorliegende Arbeit selbstständig verfasst und keine anderen als die angegebenen Quellen und Hilfsmittel verwendet habe.

Statement in Lieu of an Oath

I hereby confirm that I have written this thesis on my own and that I have not used any other media or materials than the ones referred to in this thesis.

Saarbrücken, June 21, 2025,

(Tim Christmann)

Einverständniserklärung

Ich bin damit einverstanden, dass meine (bestandene) Arbeit in beiden Versionen in die Bibliothek der Informatik aufgenommen und damit veröffentlicht wird.

Declaration of Consent

I agree to make both versions of my thesis (with a passing grade) accessible to the public by having them added to the library of the Computer Science Department.

Saarbrücken, June 21, 2025,

(Tim Christmann)

TODO: Dedication page

Yes, you can

Abstract

Trusted Web Activities and Custom Tabs enable Android developers to seamlessly integrate web content into native applications, offering a powerful tool for features such as Single Sign-On and in-app monetization. However, as shown by HyTrack, this integration also introduces severe privacy risks by blurring the boundary between web and app contexts, allowing persistent tracking through the browser’s shared cookie storage.

In this work, we propose a novel mitigation framework that applies capability-based access control to browser cookie handling. Cookie access is encapsulated in fine-grained, identity-bound capabilities, ensuring that only trusted first-party or explicitly authorized third-party web servers – defined by a developer-controlled policy – can access the shared browser state. All other untrusted third-party servers are confined to isolated, in-app cookie jars. This empowers well-meaning developers to continue leveraging third-party libraries while preventing them from performing unauthorized cross-app tracking. At the same time, essential features such as Single Sign-On and personalized content delivery remain fully functional. Our approach strikes a deliberate balance between privacy and usability, enabling secure web-app integration regarding tracking without compromising user experience.

Acknowledgements

Thanks Sven and Noah

Contents

Abstract	vii
Acknowledgements	ix
List of Figures	xiii
List of Tables	xv
1 Introduction	1
2 Related Work	3
3 Methodology: Preventing Cross-App Tracking in CTs and TWAs via Capabilities	5
3.1 Proposed Solution: Capability-Based Cookie Isolation	5
3.2 Benefits	7
3.3 Hypotheses	8
3.4 Components	9
4 Evaluation	11
5 Schedule	13
6 Success Criteria	15
6.1 Must-have criteria	15
6.2 May-have criteria	15
6.3 Must-not-have criteria	16
Bibliography	17
Appendix	19

List of Figures

List of Tables

Chapter 1

Introduction

In recent years, Android applications have increasingly leveraged web content within their interfaces to enhance user experience and streamline features such as authentication and monetization. To enable this, developers often use Custom Tabs (CTs) and Trusted Web Activities (TWAs), technologies that provide seamless, browser-backed web integration while maintaining native-like performance and features. This approach allows web-based functionality like Single Sign-On (SSO), for example login via Facebook or embedded advertising, without forcing users to switch between app and browser.

However, these benefits come at a cost. CTs and TWAs share the browser's cookie storage across all apps, enabling continuity of web sessions – but also opening serious privacy vulnerabilities. Recent research by Wessels et al. introduced HyTrack [1], a novel tracking technique that exploits this shared browser state to persistently track users across different applications and the web, even surviving device changes, cookie clearing, or browser switching. HyTrack works by embedding a third-party library into multiple unrelated apps. Each app, unaware of the library's true purpose, opens a CT or TWA to the same tracking domain. This domain sets a unique identifier in a cookie, stored in the browser's shared cookie jar. When another app using the same library loads content from the same domain, the cookie is sent, enabling the tracker to correlate activity across apps and even into regular browser use. Due to Android's backup mechanisms, the tracking ID can be restored even after a factory reset, rendering it more persistent than the evercookie [2].

This thesis explores whether capabilities, a fine-grained access control model, can be used to limit or prevent these privacy issues without breaking legitimate use cases of CTs and TWAs. Specifically, we aim to design and evaluate a framework that allows developers to retain the benefits of third-party libraries (e.g., SSO or monetization) without exposing users to invisible, cross-app tracking. The framework should be simple to integrate,

practical in real-world deployments, and minimize interference with already established app workflows.

Chapter 2

Related Work

Tracking mechanisms are typically divided into two broad categories: stateful and stateless tracking. Stateless tracking, also known as fingerprinting, infers a user’s identity based on a combination of device-specific attributes. Consequently, this method is hard to detect and block, but is also inherently less reliable, as small system changes may alter the fingerprint and disrupt identification.

Instead, stateful tracking relies on storing unique identifiers on the client device, most commonly through cookies or local storage. When a user revisits a site or interacts with embedded third-party content across domains, these identifiers are sent along with requests, allowing persistent recognition. While straightforward and highly effective, stateful tracking has become increasingly restricted through browser policies (e.g., third-party cookie blocking) and mobile platform changes such as the ability to disable the Google Advertising ID (GAID) on Android.

HyTrack [1] demonstrates a novel cross-app and cross-web tracking technique in the Android ecosystem by exploiting the shared cookie storage between Custom Tabs and Trusted Web Activities (TWAs). This allows to persistently track users across multiple applications and the browser, even surviving user efforts to reset or sanitize their environments. While HyTrack highlights a serious privacy vulnerability, no concrete mitigation has been proposed that balances privacy with the legitimate need for seamless web integration – such as Single Sign-On or ad delivery – within mobile apps.

This can be seen by taking a closer look at two mitigation strategies discussed by the authors, namely Browser State Partitioning and Forced User Interaction. However, both approaches introduce significant drawbacks. Browser state partitioning would allow each app to use its own cookie storage and hence prevent cross-app tracking. The seamless integration of web content remains intact, as no changes to the UI are necessary, but by

completely removing the browser's shared state, benign uses like Single Sign-On (SSO) or ad personalization would be broken.

In contrast, Forced User Interaction avoids this problem by allowing the browser to use its shared cookie storage. But this introduces a significant usability issue, as the user is forced to interact with the browser every time a web content is loaded, which is not only annoying but also breaks the seamless integration of web content into the app. Furthermore, this approach hands control and responsibility to the user, which is not ideal from a security perspective, as the user might be unaware of the consequences of their actions and may inadvertently enable tracking by failing to interact with the browser as required. Other strategies, such as Limiting CTs and TWAs to First-Party Domains or introducing browser options to disable CTs and TWAs for certain domains boil down to the strategies mentioned above and hence are not workable as remedies against HyTrack.

This work addresses this gap by proposing a capability-based access control framework for Android applications using CTs and TWAs. By wrapping Cookies into fine-grained capability tokens created by the browser according to the app developer's policy, the browser decides which cookies are allowed to be stored in the (by default) shared cookie storage and which are stored in the app-specific storage, and thus not accessible to third-party libraries. This ensures that there is no cross-app tracking possible for untrusted third-party libraries, as each app stores its own tracking cookie. As the shared cookie storage still exists for domains declared as first-party or trusted by the app developer, legitimate uses of the shared browser state (e.g. SSO) are preserved. Seamless integration of web content remains also untouched, as there is no need for user interaction or changes to the UI. Thus, in contrast to prior discussed mitigation strategies, this approach provides developers with a practical and enforceable way to render cross-app tracking infeasible.

Chapter 3

Methodology: Preventing Cross-App Tracking in CTs and TWAs via Capabilities

Custom Tabs (CTs) and Trusted Web Activities (TWAs) enable a seamless integration between apps and the web by sharing browser state – notably session cookies. While this improves user experience, it also introduces a significant privacy concern: Third-party libraries can exploit this shared browser state to track users across apps that embed them. HyTrack, for example, requires only a single shared third-party library used across multiple apps to persistently identify and track users, circumventing typical browser or OS sandboxing.

This solution seeks to preserve the seamless user experience of CTs/TWAs while enforcing cookie isolation across apps to prevent unauthorized cross-app tracking – particularly from embedded third-party libraries.

3.1 Proposed Solution: Capability-Based Cookie Isolation

To address the issue, we propose a developer-defined policy mechanism paired with cryptographically enforced capabilities that dictate how cookies are managed and isolated per app.

Upon app installation, the installer extracts, error checks, encrypts and finally sends such a policy to the browser. This policy – declared by the app developer in the manifest and config file – defines the following:

- A list of trusted domains (e.g., the developer's domains or selected third parties).
- A set of predefined cookie names expected to be used by those servers.

According to this policy, the browser creates capabilities for the app wrapping important cookie meta data (e.g., global jar flag, domain, name and value) into a secure structure, encoded and signed by the browser similar to JSON Web Tokens (JWT) [3]. Consequently, these capabilities act as **authorization tokens** for cookie access and fall into one of the following categories:

- **Wildcard Capabilities:** Allow a webserver to define arbitrary cookie names and values by leaving these fields unspecified in the capability, which are later filled by the browser when it receives a cookie from the webserver. This mechanism effectively transforms regular cookies into structured capabilities wrapping the cookie. Hence, the app receives such a capability for each domain registered in the policy.
- **Predefined Capabilities:** Define fixed cookie names as specified in the policy and can be stored either in the app-specific cookie jar or the browser's shared cookie jar, depending on the configuration of the global jar flag. This provided fine-grained control over cookie scoping and allows the developers to isolate their app from the browser's state when desired, while still enabling seamless web-app integration.
- **Ambient Capabilities:** Instruct the browser to revert to default behavior by storing all received cookies in the global cookie jar and including them in subsequent requests to the corresponding web server. These are created by setting the global jar flag and extending a wildcard capability to include the domain name, cookie name, and value.

Ambient capabilities do not provide security guarantees and exist solely as a fallback to maintain program flow when no explicit policy is defined by the developer.

When the app opens a URL via CT or TWA, it includes a list of capabilities along with the regular intent with the target URL.

The browser parses and verifies each capability by checking the following fields:

- **Signature:** Ensures the capability was issued by the browser and has not been tampered with.

- **App ID:** Identifies the origin of the request, ensuring that only authorized apps can hold a given capability. This field is essential to prevent implicit delegation – without it, a receiving app could exploit the capabilities granted to the sending app and potentially gain unauthorized access.
- **Domain:** Specifies the intended destination web server for the capability, ensuring it is sent to the correct endpoint.
- **App Version Number:** Indicates the version of the app to detect potential policy changes and ensure consistency between app and policy.
- **Rights:** Define the scope of access granted to the app – for example, whether it may request the browser to read cookies. This restriction is essential to prevent libraries from exploiting browser access to extract capability values, which could otherwise be used for tracking.
- **Global Jar Flag:** Determines whether the shared cookie jar or an app-specific cookie jar should be used.

The communication between Browser and Webserver remains unchanged, i.e. the browser sends Request and Set-Cookie headers as usual and the webserver responds with an (customized) Response and possibly new Cookies.

The browser then matches received cookies to capabilities:

- If the cookie name is predefined, only the value is updated in the capability.
- If it's new and a wildcard capability exists, the browser creates a new capability by copying the wildcard capability, setting the name and value accordingly and storing it in the browser's cookie jar. The initial wildcard capability is then sent back to the app.
- If it's new and no wildcard capability exists, the browser creates a new capability with the name and value, domain and app ID, restricted rights and version number, and hands it back to the app to store it in its app-specific cookie jar.

3.2 Benefits

Beyond eliminating cross-app tracking through HyTrack and their postulated goals, we identify the following additional benefits offered by this approach:

- B1) **Fine-Grained Control:** Developers can specify which cookies are shared and which are isolated, allowing for a more tailored approach to privacy.
- B2) **No Browser State for Apps:** The browser does not need to remember app capabilities – the app holds and re-sends them with each intent.
- B3) **No Third-Party Code Changes:** The webserver does not need to be aware of the capabilities or make any changes to its code, as the browser handles the capability management.
- B4) **Backwards Compatibility:** In case the developer does not implement its own policy, it will default to allowing any domain using the shared browser state, consequently behaving like the state-of-the art cookie management.

3.3 Hypotheses

The evaluation is based on the following hypotheses:

- H1) Passing capabilities with each intent enables the browser to validate and manage cookies securely and effectively. It also ensures the browser does not need to store app states.
- H2) The installer can transmit developer-defined policies to the browser and return them to the requesting app reliably.
- H3) The proposed mechanism is compatible with the original three goals of HyTrack as outlined in Section 4.
- H4) The mechanism significantly mitigates cross-app tracking by isolating browser state based on trusted policy definitions.
- H5) Backwards compatibility with existing systems can be ensured through the use of *Ambient Capabilities*.
- H6) In the case of invalid or missing capabilities, a default-deny policy can be enforced via empty requests, preventing fallback to shared state.
- H7) The proposed approach requires no modifications to existing third-party webserver code.

3.4 Components

In summary, the proposed solution requires the following components to be modified or extended:

- **Manifest** and **config file** for the developers to define which origins should retain shared browser state.
- **Custom Installer** to extract, error check, encrypt and send the developer-defined policy to the browser upon app installation.
- **Browser** to function as Policy Enforcement Point. For this, new functionality for creating capabilities according to the policy and validating them is necessary. In addition, it needs to be capable of storing a secret key for signing capabilities and verifying the signature. Ofcourse, decision making according to the capabilities context is unavoidable, i.e. where to store the capability, or the pair of name and value, either in the app-specific or the global cookie jar.
- **App** to send capabilities with each intent and to store the capabilities returned by the browser. It must also support functionality to send specific capabilities to the browser.
- **Capabilities** to wrap the regular cookie data, app ID, domain, global Jar flag and rights.

Chapter 4

Evaluation

To assess the effectiveness of our proposed mitigation strategy, we adopt the three primary goals identified by the authors of HyTrack as essential for any viable defense:

- 1) **Support all features of the web platform:** The solution must allow applications to display fully functional web content, including support for cookies, JavaScript, and modern APIs.
- 2) **Preserve seamless integration:** The user experience must remain uninterrupted. This includes avoiding obtrusive permission dialogs and maintaining smooth transitions between native and web content.
- 3) **Enable controlled access to shared browser state:** While isolation is required to prevent cross-app tracking, legitimate scenarios such as Single Sign-On (SSO) must remain functional.

These criteria reflect the fact that HyTrack exploits standard Android behavior – specifically, the shared browser state exposed through Custom Tabs and Trusted Web Activities – rather than relying on unauthorized access or system vulnerabilities. Therefore, naive approaches like disabling shared cookies entirely would break common use cases and are not acceptable.

To validate these hypotheses, we will build on the open-sourced measurement tooling and proof of concept applications provided by the authors of HyTrack. Specifically, we plan to:

- Replicate the original HyTrack experiments under controlled conditions using two unrelated Android apps that embed the tracking library (similar to the HyTrack demo).

- Apply the mitigation framework and compare observed behavior against the baseline.

We will collect and analyze the following metric:

- Number of Capabilities created and used by the browser.

In doing so, we aim to demonstrate that the proposed solution effectively blocks HyTrack's cross-app tracking channel while preserving compatibility with existing web features and maintaining seamless user experience. Additionally, we will show that this strategy can be used by developers to control cookie transmission in a fine-grained and policy-driven manner to enable safer and more transparent integration of third-party web content.

Chapter 5

Schedule

We plan to complete the project in the following phases:

- **Week 1: Background Research and Literature Review**

Study existing work on web and cross-app tracking, especially the HyTrack paper. Familiarize with Android's Custom Tabs (CTs), Trusted Web Activities (TWAs), and Chromium's cookie storage mechanisms. Reproduce the HyTrack proof of concept to establish a baseline for evaluation.

- **Weeks 2-4: Implementation**

Identify which components of Chromium and App need to be modified. Implement the capability-based cookie isolation framework.

- **Week 5: Testing**

Thoroughly test the framework to ensure it works as intended and meets the requirements.

- **Week 6: Evaluation and Experimentation**

Set up experimental infrastructure using HyTrack's open-source tooling. Measure the effectiveness of the proposed solution by replicating tracking scenarios and comparing cookie behaviors across test cases. Conduct experiments to evaluate the framework's performance, feasibility of integration and backwards compatibility.

Risks, Impact and Mitigation

At worst, the capability-based approach may not effectively prevent cross-app tracking in practice and indeed a complete redesign of the Custom Tab and Trusted Web Activity

API is needed, as suggested by the authors of HyTrack. If this occurs, the thesis will pivot to a critical analysis of why the capability model falls short in this context. The focus will shift towards identifying structural barriers in the Android platform and recommending future improvements to make such defenses feasible.

In case we encounter problems in adapting the Installer to send the policy to the browser, we will explore sending the policy on each launch of the app via a modified Zygote instead.

At much lower impact, the implementation may not be feasible within the given timeframe or due to unforeseen technical challenges. In this case, development will focus on a minimally viable proof-of-concept that demonstrates core functionality and missing features will be discussed as future work.

Chapter 6

Success Criteria

The following criteria define the scope of this thesis and provide measurable goals to assess its success.

6.1 Must-have criteria

- The solution must effectively prevent or significantly limit cross-app tracking via the shared browser state, while preserving the core functionality of Custom Tabs and Trusted Web Activities (e.g., Single Sign-On).
- It must adhere to the three primary goals outlined by HyTrack:
 - (1) Support all features of the Web platform,
 - (2) do not break the seamless integration, and
 - (3) make shared state available to the Custom Tab or Trusted Web Activity.
- A working proof-of-concept must be implemented to demonstrate the feasibility of the approach and to reproduce and compare against the HyTrack attack methodology.
- The thesis must provide a thorough discussion of how capability-based access control contributes to mitigating the identified tracking risks, including analysis of potential trade-offs (e.g., small form of ambient authority as based on identity).

6.2 May-have criteria

- Provide an analysis of the solution's impact on attributes like performance, usability, integration effort for developers, and potential security limitations.

- Evaluate the framework across multiple Chromium-based browsers to assess generalizability.
- Investigate whether enhancements to the Digital Asset Links (DAL) mechanism could further strengthen the solution or propose a good alternative.

6.3 Must-not-have criteria

- The thesis will not produce a production-ready implementation, as the focus is on proof-of-concept and feasibility.
- The solution is not intended to prevent all forms of cross-app tracking, particularly not against malicious developers who intentionally bypass protections.
- It will not attempt to mitigate stateless tracking methods such as fingerprinting, which are outside the scope of shared state via cookies.

Bibliography

- [1] M. Wessels, S. Koch, J. Drescher, L. Bettels, D. Klein, and M. Johns, “Hytrack: Resurrectable and persistent tracking across android apps and the web,” in *34th USENIX Security Symposium (USENIX Security 25)*. Seattle, WA: USENIX Association, Aug. 2025.
- [2] S. Kamkar, “evercookie - virtually irrevocable persistent cookies,” <https://samy.pl/evercookie/>, 2010, accessed: 2025-05-25.
- [3] M. B. Jones, J. Bradley, and N. Sakimura, “JSON Web Token (JWT),” RFC 7519, May 2015. [Online]. Available: <https://www.rfc-editor.org/info/rfc7519>

Appendix

Notes

■ TODO: Dedication page v